

# CSI2110

## Data Structures and Algorithms

- <https://www.youtube.com/watch?v=aQS9DqLWxw4>
- <https://www.cs.usfca.edu/~galles/visualization/AVLtree.html>


# AVL Trees

Adel'son-Vel'skii and Landis

---

Data structure that implements MAP ADT

- Height of an AVL Tree
- Insertion and restructuring
- Removal and restructuring
- Costs




# AVL Tree

---

- AVL trees are balanced.
- An AVL Tree is a *binary search tree* such that for every internal node  $v$  of  $T$ , the *heights of the children of  $v$*  can differ by at most 1.


An example of an AVL tree where the heights are shown next to the nodes:



# Balancing Factor

$\text{height}(\text{right subtree}) - \text{height}(\text{left subtree})$

$\in \{-1, 0, 1\}$  for AVL tree



# Height of an AVL Tree

---


Note: “longest” possible heap with  $n$  nodes.

Always  $O(\log n)$

But cannot search efficiently

Note: “longest” possible binary tree with  $n$  nodes:


$O(n)$



Inefficient to search in the worst case.

We'll now see that the *height* of an AVL tree  $T$  storing  $n$  keys is  $O(\log n)$ .



*Note:* AVL tree with the highest possible number of internal nodes for a given height  $h$ :




To construct the “longest” possible AVL tree, we look for the *minimum number of nodes of an AVL tree of height  $h$ .*  $n(h)$

Easy to see that  $n(1) = 1$  and  $n(2) = 2$

(please note that dummy nodes are not shown but contributing to the height)



$$n(1) = 1$$




2

$$n(2) = 2$$



$n(h)$ : the *minimum number of internal nodes* of an AVL tree of height  $h$ .

For  $n \geq 3$ , an AVL tree of height  $h$  contains the root node, one AVL subtree of height  $h-1$  and the other AVL subtree of height  $h-2$ .




$$n(h) = 1 + n(h-1) + n(h-2)$$


# Height of an AVL Tree

---

$$n(h) = 1 + \boxed{n(h-1)} + n(h-2)$$



↑  
h-1  
↓



h-2  
↑  
↓



Clearly:

$$\boxed{n(h-1) > n(h-2)}$$

so:  $n(h) > 1 + n(h-2) + n(h-2)$   
 $n(h) > 1 + 2n(h-2) > 2n(h-2)$



$$n(h) > 2n(h-2)$$

for any h

## Height of an AVL Tree

---

So, now we know:  $n(h) > 2 \boxed{n(h-2)}$   
but then also:  $\boxed{n(h-2)} > 2 n(h-4)$  }  $n(h) > 4 n(h-4)$

$n(h) > \boxed{4n(h-4)}$   
but then also:  $\boxed{n(h-4)} > 2n(h-6)$  }  $n(h) > 8n(h-6)$

We can continue:

$$n(h) > 2n(h-2)$$

$$n(h) > 4n(h-4)$$

$$n(h) > 8n(h-6)$$

...

$$n(h) > 2^i n(h-2i)$$

$$n(h) > 2^i n(h-2i)$$



$$h-2i = 2$$

And we know that

$$n(1) = 1$$

$$n(2) = 2$$

Now we pick  $i$  such that  $h$  is either 1 or 2

That is pick  $i = \text{ceil}(h/2) - 1$ , and substitute:

$$\text{for } i = \text{ceil}(h/2) - 1 \quad n(h) > 2^{\text{ceil}(h/2) - 1} n(1)$$

$$n(h) > 2^{h/2-1}$$

$$\log n(h) > \log 2^{(h/2)-1}$$

$$\log n(h) > h/2 - 1$$

$$h < 2 \log n(h) + 2$$

which means that  $h$  is  $O(\log n)$


# Insertion

---


- A binary search tree  $T$  is called *balanced* if for every node  $v$ , the height of  $v$ 's children differ by at most one.
- Inserting a node into an AVL tree involves performing an *expandExternal( $w$ )* on  $T$ , which changes the heights of some of the nodes in  $T$ .
- If an insertion causes  $T$  to become *unbalanced* we have to rebalance...

# Insertion








Before



After left  
insertion



After right  
insertion



# Rebalancing after insertion



We are going to identify 3 nodes which form a grandparent, parent, child triplet and the 4 subtrees attached to them. We will rearrange these elements to create a new balanced tree.



# Rebalancing

**Step 1:** Trace the path back from the point of insertion to the first node whose **grandparent is unbalanced**. Label this node  $x$ , its parent  $y$ , and grandparent  $z$ .



Examples .....



# Rebalancing

**Step 1:** Trace the path back from the point of insertion to the first node whose **grandparent is unbalanced**. Label this node  $x$ , its parent  $y$ , and grandparent  $z$ .



Examples .....



# Rebalancing

**Step 2:** These nodes will have 4 subtrees connected to them. Label them  $T_1$ ,  $T_2$ ,  $T_3$ ,  $T_4$  from left to right.


Examples .....



# Rebalancing


**Step 3:** Rename  $x, y, z$  to  $a, b, c$  according to their inorder traversal i.e. if  $y, x, z$  is the relative order of those nodes following the inorder traversal then label  $y$  'a',  $x$  'b' and  $z$  'c'.

## Example



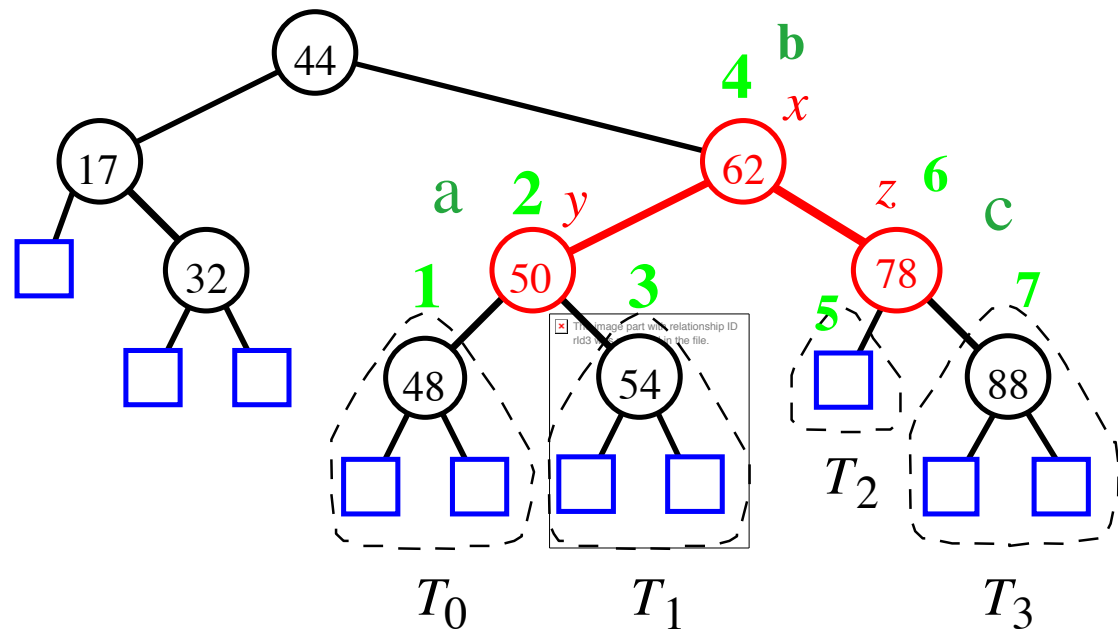
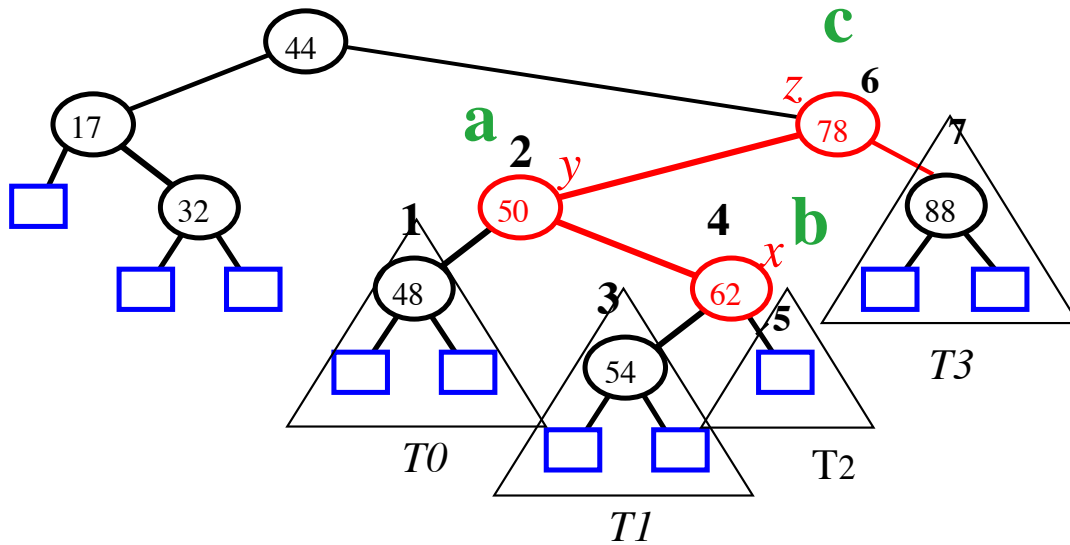
# Rebalancing

**Step 4:** Replace the tree rooted at  $z$  with the following tree:



Rebalance done!

Example: after  
inserting 54



# Does this really work?

We need to see that the new tree is :



- a) A Binary search tree - the inorder traversal of our new tree should be the same as that of the old tree

Inorder traversal: by definition is T1 a T2 b T3 c T4

- b) Balanced: have we fixed the problem?

We consider 2 types of examples



# Example 1



Inorder:  $T_1 z T_2 y T_3 x T_4$



# Example 2



Inorder: T1 y T2 x T3 z T4

# An Observation...

Notice that in both cases, the new tree rooted at  $b$  has the same height that the old tree rooted at  $z$  had before insertion.

So.. once we have done one rebalancing act, we are done.


## rebalance (v)

```
x <- v; Y <- x.parent; z <- y.parent
while (z.isBalanced and not(z.isRoot) )
  x <- y; y <- z; z <- z.parent
if ( not z.isBalanced)
  if ( x = y.left) { x<=y}
    if (y = z.left) {x<=y<=z}
      a <- x; b <- y; c <- z;
      T2 <- x.right; T3 <- y.right;
    else { z<=x<=y}
      a <- z; b <- x; c <- y;
      T2 <- x.left; T3 <- x.right;
  else {y<=x}
    if (y = z.left) {y<=x<=z}
      a <- y; b <- x; c <- z;
      T2 <- x.left; T3 <- x.right
    else { z<=y<=x}
      a <- z; b <- y; c <- x;
      T2 <- y.left; T3 <- x.left
```


```
T1 <- a.left; T4 <- c.right
b.left <- a; b.right <- c
a.left <- T1; a.right <- T2
c.left <- T3; c.right <- T4
T1.parent <- a; T2.parent <- a
T3.parent <- b; T4.parent <- c

if (z.isRoot) then
  root <- b
  b.parent <- NULL
else if (z.isLeftChild)
  z.parent.left <- b
  else z.parent.right <- b
b.parent <- z.parent
a.parent <- b; c.parent <- b
```

# Restructuring (as Single Rotations)



# Restructuring (as Double Rotations)



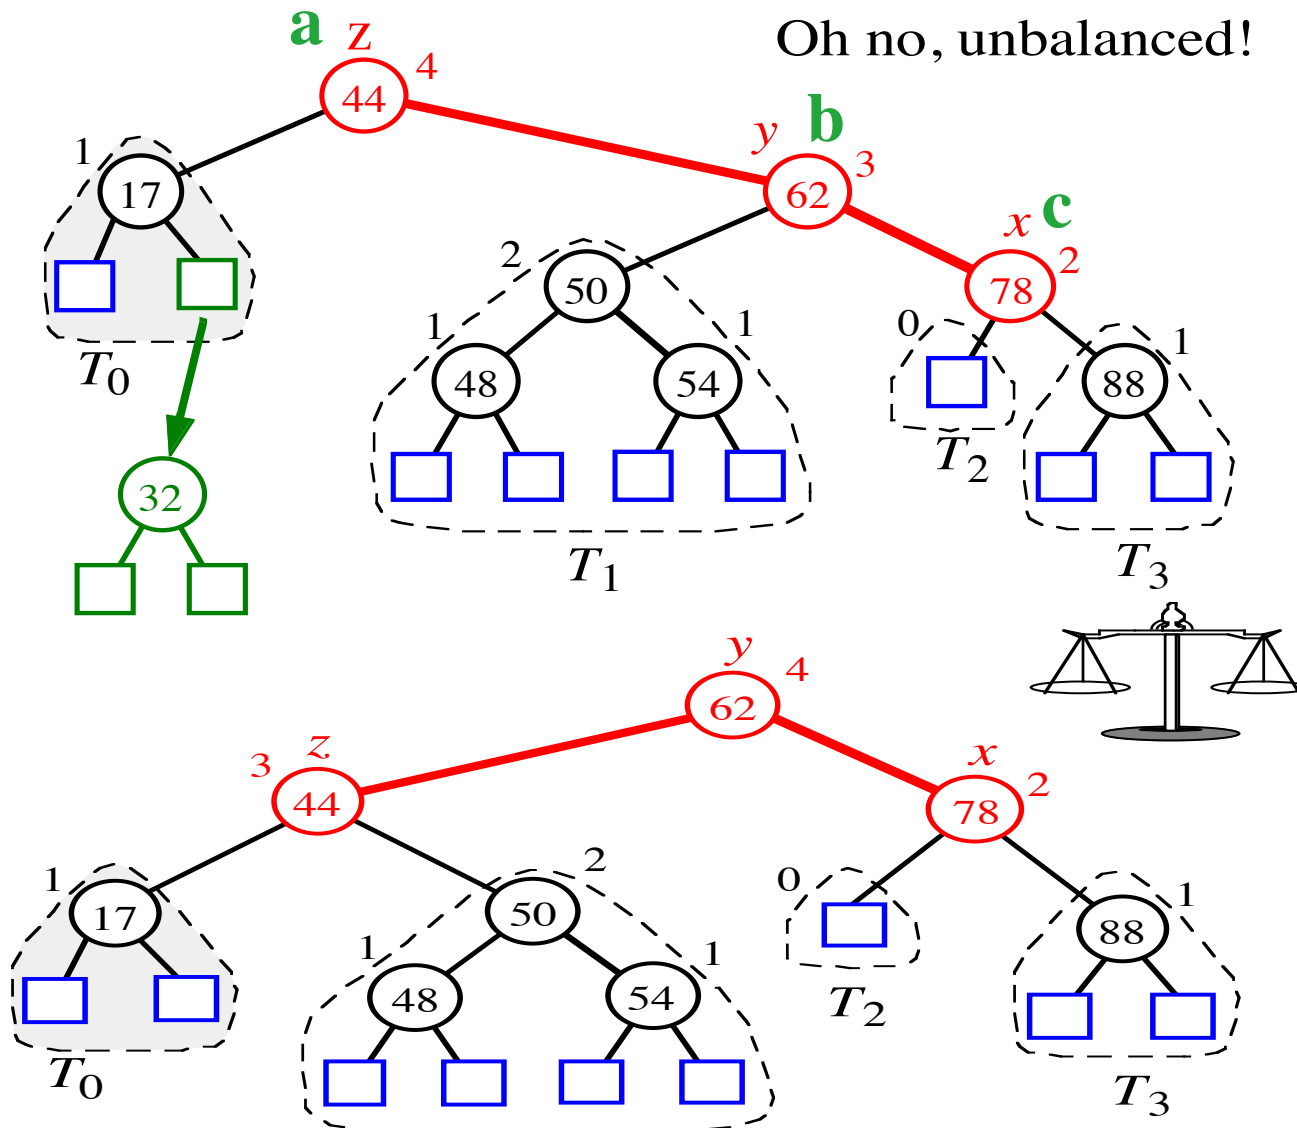
# Removal

---

- We can easily see that performing a `removeAboveExternal(w)` can cause  $T$  to become unbalanced.
- Let  $z$  be the **first unbalanced** node encountered while travelling up the tree from  $w$ . Also, let  $y$  be **the child of  $z$  with the larger height**, and let  $x$  be **the child of  $y$  with the larger height**.
- We can perform operation `restructure(x)` to restore balance at the subtree rooted at  $z$ .
- As this restructuring may upset the balance of another node higher in the tree, we must continue checking for balance until the root of  $T$  is reached

# Removal (contd.)


the choice of  $x$  is not unique !!!




Whew,  
balanced!

# Removal (contd.)

- we could choose a different  $x$ :



Oh no,  
unbalanced!



Whew,  
balanced!



Again..  $x, y, z$

$z$  : the first unbalanced node encountered while travelling up the tree from  $w$ .

$y$  : the child of  $z$  with the larger height

$x$  : the child of  $y$  with the larger height. If both children of  $y$  have the same height, let  $x$  be the child of  $y$  on the same side as  $y$ .


# COMPLEXITY

Searching:	findElement(k):
Inserting:	insertItem(k, o):
Removing:	removeElement(k):

$O(\log n)$

Some implementation details are very important:

The trinode restructure is accomplished using the rotation operation:

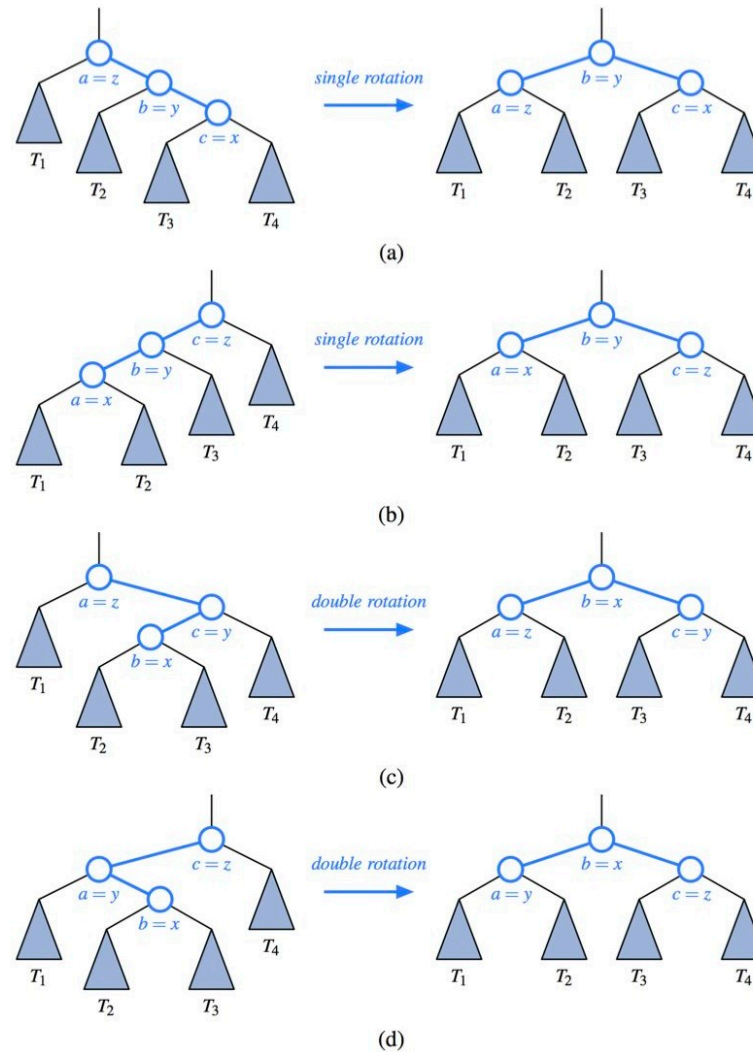


**Figure 11.8:** A rotation operation in a binary search tree. A rotation can be performed to transform the left formation into the right, or the right formation into the left. Note that all keys in subtree  $T_1$  have keys less than that of position  $x$ , all keys in subtree  $T_2$  have keys that are between those of positions  $x$  and  $y$ , and all keys in subtree  $T_3$  have keys that are greater than that of position  $y$ .

# Trinode restructuring using rotation operation:

474

Chapter 11. Search Trees



**Figure 11.9:** Schematic illustration of a trinode restructuring operation: (a and b) require a single rotation; (c and d) require a double rotation.

Rotate:

```
28  /** Relinks a parent node with its oriented child node. */
29  private void relink(Node<Entry<K,V>> parent, Node<Entry<K,V>> child,
30      boolean makeLeftChild) {
31      child.setParent(parent);
32      if (makeLeftChild)
33          parent.setLeft(child);
34      else
35          parent.setRight(child);
36  }
37  /** Rotates Position p above its parent. */
38  public void rotate(Position<Entry<K,V>> p) {
39      Node<Entry<K,V>> x = validate(p);
40      Node<Entry<K,V>> y = x.getParent();           // we assume this exists
41      Node<Entry<K,V>> z = y.getParent();           // grandparent (possibly null)
42      if (z == null) {
43          root = x;                               // x becomes root of the tree
44          x.setParent(null);
45      } else
46          relink(z, x, y == z.getLeft());           // x becomes direct child of z
47      // now rotate x and y, including transfer of middle subtree
48      if (x == y.getLeft()) {
49          relink(y, x.getRight(), true);             // x's right child becomes y's left
50          relink(x, y, false);                       // y becomes x's right child
51      } else {
52          relink(y, x.getLeft(), false);             // x's left child becomes y's right
53          relink(x, y, true);                       // y becomes left child of x
54      }
55  }
56  /** Performs a trinode restructuring of Position x with its parent/grandparent. */
57  public Position<Entry<K,V>> restructure(Position<Entry<K,V>> x) {
58      Position<Entry<K,V>> y = parent(x);
59      Position<Entry<K,V>> z = parent(y);
60      if ((x == right(y)) == (y == right(z))) {     // matching alignments
61          rotate(y);                                // single rotation (of y)
62          return y;                                 // y is new subtree root
63      } else {                                       // opposite alignments
64          rotate(x);                                // double rotation (of x)
65          rotate(x);
66          return x;                                 // x is new subtree root
67      }
68  }
69  }
```

Restructure:

**Code Fragment 11.10:** The `BalanceableBinaryTree` class, which is nested within the `TreeMap` class definition (continued from Code Fragment 11.9).

# Rebalancing operation for AVL insertions and deletions:

```
28  /**
29   * Utility used to rebalance after an insert or removal operation. This traverses the
30   * path upward from p, performing a trinode restructuring when imbalance is found,
31   * continuing until balance is restored.
32   */
33  protected void rebalance(Position<Entry<K,V>> p) {
34      int oldHeight, newHeight;
35      do {
36          oldHeight = height(p);                // not yet recalculated if internal
37          if (!isBalanced(p)) {                  // imbalance detected
38              // perform trinode restructuring, setting p to resulting root,
39              // and recompute new local heights after the restructuring
40              p = restructure(tallerChild(tallerChild(p)));
41              recomputeHeight(left(p));
42              recomputeHeight(right(p));
43          }
44          recomputeHeight(p);
45          newHeight = height(p);
46          p = parent(p);
47      } while (oldHeight != newHeight && p != null);
48  }
49  /** Overrides the TreeMap rebalancing hook that is called after an insertion. */
50  protected void rebalanceInsert(Position<Entry<K,V>> p) {
51      rebalance(p);
52  }
53  /** Overrides the TreeMap rebalancing hook that is called after a deletion. */
54  protected void rebalanceDelete(Position<Entry<K,V>> p) {
55      if (!isRoot(p))
56          rebalance(parent(p));
57  }
58  }
```



At this point in the class I discuss how AVL trees are implemented in the 6<sup>th</sup> edition of the textbook by Goodrich, Tamassia and Goldwasser.

Please, refer to pages:

466-470 class `TreeMap<K,V>`

Note methods: `put(K key, V value)`, `remove(K key)`

475-478 class `BalancedBinaryTree<K,V>`

Note: hooks for rebalancing present in `TreeMap`,

Methods: `rotate`, `restructure`

486-487 class `AVLTreeMap<K,V>`

Note methods: `rebalanceInsert`, `rebalanceDelete`, `rebalance`